

pyLGC — Python Interface for LGC2

Least-Squares Geodetic Computation

Jürgen Gutekunst, Marco Carotta

June 18, 2026

Contents

1	Introduction	2
2	The Least-Squares Problem	2
2.1	The Nonlinear Estimation Problem	2
2.2	Gauss–Newton Linearisation	2
2.3	Solution of the Linearised Problem	3
3	API Reference	3
3.1	Architecture	3
3.2	Module-Level Types	4
3.3	Class <code>Evaluator</code>	4
3.4	Class <code>AdjustableObject</code>	6
3.5	Error Handling	6
4	Typical Workflow and Examples	6
4.1	Sparse Matrix Usage with SciPy	8
4.2	Building the Normal Equations in SciPy	8
4.3	Gauss–Newton Iteration	9

1 Introduction

`pyLGC` is a Python module that provides access to the LGC2 least-squares adjustment engine. It loads a native shared library (`pyLGC_C.dll` on Windows, `libpyLGC_C.so` on Linux) via `ctypes`.

The module exposes the evaluator, which encapsulates the linearised least-squares problem built from an LGC input file, and allows the user to inspect, perturb, evaluate, and solve the underlying system. For a detailed description of the underlying concepts and methods, see [1].

This document is structured as follows: Section 2 describes the underlying least-squares problem, Section 3 documents the Python API, and Section 4 provides examples of common workflows.

2 The Least-Squares Problem

This section summarises the nonlinear estimation problem solved by LGC2. The formulation follows [2]; see also [1] for a broader overview and [3] for the general theory of least squares.

In the general theory of geodetic adjustment, the observation model can take the implicit (combined) form $F(x, L + v) = 0$, where both parameters and observations appear inside F . In LGC2, however, all observation models are of *parametric* form: $f(x) - (L + v) = 0$, i.e. the model predicts the observations as a function of the parameters alone. As a consequence, the matrix $B = \partial F / \partial L$ is always $-I$ (the negative identity), which simplifies the solution considerably.

2.1 The Nonlinear Estimation Problem

Given a parametric observation model $f: \mathbb{R}^{n_x} \rightarrow \mathbb{R}^{n_L}$ that predicts measurement values from unknown parameters, LGC2 solves the constrained nonlinear least-squares problem

$$\begin{aligned} & \min \frac{1}{2} v^T P v \\ & \text{over } (x, v) \in \mathbb{R}^{n_x} \times \mathbb{R}^{n_L}, \\ & \text{subject to} \\ & \quad f(x) + B v - L = 0, \\ & \quad C(x) = 0, \end{aligned} \tag{1}$$

where $x \in \mathbb{R}^{n_x}$ is the parameter vector, $L \in \mathbb{R}^{n_L}$ the measurement vector, $v \in \mathbb{R}^{n_L}$ the residuals, $B \in \mathbb{R}^{n_L \times n_L}$ the observation coefficient matrix, $C: \mathbb{R}^{n_x} \rightarrow \mathbb{R}^{n_C}$ the constraint function, and $P \in \mathbb{R}^{n_L \times n_L}$ the weight matrix. Since all observation models in LGC2 are of parametric form $f(x) - (L + v) = 0$, the matrix B is always $-I$ (negative identity). The API nevertheless exposes B and B^{-1} for completeness.

2.2 Gauss–Newton Linearisation

Problem (1) is solved by the Gauss–Newton algorithm. At iterate x_k , we define the misclosure vector as

$$W_k = f(x_k) - L \in \mathbb{R}^{n_L}, \tag{2}$$

which measures how well the current parameters explain the measurements. Linearising f around x_k gives

$$f(x_k + \Delta x) \approx f(x_k) + A_k \Delta x, \quad A_k = \frac{\partial f}{\partial x}(x_k) \in \mathbb{R}^{n_L \times n_x}. \tag{3}$$

Similarly, the constraint is linearised as

$$C(x_k + \Delta x) \approx C(x_k) + A_{2,k} \Delta x, \quad A_{2,k} = \frac{\partial C}{\partial x}(x_k) \in \mathbb{R}^{n_C \times n_x}. \tag{4}$$

Substituting the linearised model into the observation equation $f(x) + Bv - L = 0$ gives

$$f(x_k) + A_k \Delta x + Bv - L \approx 0 \quad \implies \quad v \approx -B^{-1}(f(x_k) - L + A_k \Delta x) = -B^{-1}(W_k + A_k \Delta x). \quad (5)$$

With $B = -I$ this simplifies to $v \approx W_k + A_k \Delta x$. Inserting into the objective $\min \frac{1}{2} v^T P v$ yields the linearised subproblem

$$\begin{aligned} \min \quad & \frac{1}{2} (W_k + A_k \Delta x)^T P (W_k + A_k \Delta x) \quad \text{over } \Delta x, \\ \text{s.t.} \quad & A_{2,k} \Delta x + C(x_k) = 0. \end{aligned} \quad (6)$$

This defines the matrices and vectors accessible through the `pyLGC` API:

$$A = \frac{\partial f}{\partial x}(x_k) \in \mathbb{R}^{n_L \times n_x}, \quad (\text{design / Jacobian matrix}) \quad (7)$$

$$B = -I \in \mathbb{R}^{n_L \times n_L}, \quad (\text{observation coefficient matrix}) \quad (8)$$

$$W = f(x_k) - L \in \mathbb{R}^{n_L}, \quad (\text{misclosure vector}) \quad (9)$$

$$A_2 = \frac{\partial C}{\partial x}(x_k) \in \mathbb{R}^{n_C \times n_x}, \quad (\text{constraint Jacobian}) \quad (10)$$

$$W_2 = C(x_k) \in \mathbb{R}^{n_C}, \quad (\text{constraint misclosure}) \quad (11)$$

$$P \in \mathbb{R}^{n_L \times n_L}. \quad (\text{weight matrix}) \quad (12)$$

2.3 Solution of the Linearised Problem

Setting the gradient of the Lagrangian of (6) to zero yields the KKT system

$$\begin{pmatrix} A^T P A & A_2^T \\ A_2 & 0 \end{pmatrix} \begin{pmatrix} \Delta x \\ \lambda \end{pmatrix} = - \begin{pmatrix} A^T P W \\ W_2 \end{pmatrix}, \quad (13)$$

where λ is the vector of Lagrange multipliers for the constraints. In the unconstrained case ($n_C = 0$), this simplifies to the normal equations

$$A^T P A \Delta x = -A^T P W. \quad (14)$$

The parameters are then updated as $x_{k+1} = x_k + \Delta x$ and the process repeats until convergence. The residuals at convergence are $v = f(x^*) - L = W^*$.

3 API Reference

3.1 Architecture

`pyLGC` is structured in two layers:

1. `pyLGC_C.dll / libpyLGC_C.so` — A shared library with a plain C API (`extern "C"`) wrapping the C++ LGC2 engine. It has no dependency on Python and can be called from any language with a C foreign-function interface.
2. `pyLGC.py` — A pure-Python module that loads the shared library via `ctypes` and exposes the `Evaluator` and `AdjustableObject` classes (`AdjustablePoint` and `AdjustableFrame` are retained as aliases for `AdjustableObject`). Arrays are returned as `numpy` arrays for efficient data transfer. Requires `numpy`; `scipy` is optional (for sparse matrix construction).

Because the shared library has no compile-time dependency on Python headers, it can be used from any language with a C foreign-function interface.

3.2 Module-Level Types

UE0Indices Named tuple holding the four key dimensions of the problem:

Symbol	Field	Meaning
n_x	UIndex	Number of unknowns (parameters)
n_L	EIndex	Number of equations (= number of observations)
n_L	OIndex	Number of observations
n_C	CIndex	Number of constraints

3.3 Class Evaluator

The central class. Each instance encapsulates one least-squares problem loaded from an LGC input file.

Constructor

Evaluator(filePath: str)

Load and parse the given LGC input file. Raises **RuntimeError** if the file cannot be opened or contains errors.

Problem Dimensions

getProblemDimensions() -> **UE0Indices**

Return the dimensions (n_x, n_L, n_L, n_C) of the problem. Available immediately after construction.

Parameter Management

getEstimatedParameters() -> **numpy.ndarray**

Return the current parameter vector $x \in \mathbb{R}^{n_x}$ as a one-dimensional numpy array.

setParameters(para)

Set the parameter vector. Accepts any sequence of length n_x . Raises **RuntimeError** if the dimension is wrong. *Note:* after calling **setParameters**, a new **evaluate()** is required before the getters return valid results.

Evaluation

evaluate() -> **bool**

Linearise the functional model at the current parameter values. Computes all matrices and vectors. Returns **True** on success. Must be called before any matrix or vector getter.

Vectors

getMisclosure() -> **numpy.ndarray**

Return the misclosure vector $W = f(x_k) - L \in \mathbb{R}^{n_L}$. Requires a preceding **evaluate()**.

getConstraintMisclosure() -> **numpy.ndarray**

Return the constraint misclosure $W_2 = C(x_k) \in \mathbb{R}^{n_C}$. Requires a preceding **evaluate()**.

Sparse Matrices

All sparse matrix getters return a tuple of five elements:

`(rows, cols, vals, nrows, ncols)`

where `rows`, `cols` are integer arrays, `vals` is a float array (COO format), and `nrows`, `ncols` are the matrix dimensions. See Section 4.1 for how to construct `scipy.sparse` matrices from these tuples.

`getFirstDesignMatrix()`

Design matrix (Jacobian) $A = \partial f / \partial x \in \mathbb{R}^{n_L \times n_x}$.

`getSecondDesignMatrix()`

Observation coefficient matrix $B \in \mathbb{R}^{n_L \times n_L}$. Since LGC2 uses parametric models, $B = -I$ always (see Section 2).

`getConstraintDesignMatrix()`

Constraint Jacobian $A_2 = \partial C / \partial x \in \mathbb{R}^{n_C \times n_x}$.

`getWeightMatrix()`

Weight matrix $P \in \mathbb{R}^{n_L \times n_L}$.

All require a preceding `evaluate()`.

Solving

`solve() -> (bool, numpy.ndarray)`

Run the full iterative least-squares solution. Returns a tuple (`ok`, `solution`) where `ok` indicates convergence and `solution` $\in \mathbb{R}^{n_x}$ is the estimated parameter vector. Can be called without a prior `evaluate()`. *Note:* after solving, `evaluate()` must be called again before the matrix/vector getters are valid.

Observation Mapping

`getObsIndexToLineNumber() -> dict`

Return a dictionary mapping internal observation indices to line numbers in the input file. Useful for tracing residuals back to specific observations.

Point and Frame Access

`getPoint(name: str) -> AdjustableObject`

Look up an adjustable point by name. Raises `RuntimeError` if not found.

`getFrame(name: str) -> AdjustableObject`

Look up an adjustable reference frame by name. Raises `RuntimeError` if not found.

The two methods perform distinct lookups in the underlying engine (against `LGCAdjustablePoint` and `TAdjustableHelmertTransformation` respectively), but both return the same Python wrapper class (see Section 3.4).

3.4 Class AdjustableObject

Represents an adjustable object in the network — either a point or a reference frame. Instances are obtained via `Evaluator.getPoint()` or `Evaluator.getFrame()` and are valid as long as the parent `Evaluator` exists.

`AdjustablePoint` and `AdjustableFrame` are retained as aliases for `AdjustableObject`, so existing code that references those names continues to work.

`getName()` -> str

Return the object's name.

`getFirstUidx()` -> int

Return the index of this object's first unknown in the global parameter vector. Together with `getRelativeUnknIndices()`, this identifies which elements of x belong to this object.

`getRelativeUnknIndices()` -> numpy.ndarray

Return the relative unknown indices (offsets from `getFirstUidx()`) as a 1-D integer array. For a fully free 3D point this is `[0, 1, 2]`; for a partially fixed point some components are absent; for a fully fixed object the array is empty. For an adjustable Helmert frame the indices select among up to seven parameters (3 translations, 3 rotations, 1 scale).

`getEstVector()` -> numpy.ndarray

Return the current estimated parameter vector as a 1-D float array. For a 3D point this is always 3 elements regardless of which components are fixed; for a Helmert frame this is the current transformation parameter vector.

3.5 Error Handling

All errors are reported as Python `RuntimeError` exceptions with descriptive messages. Common error conditions:

Situation	Example message
File not found	Cannot open file: network.lgc2
Malformed input file	Errors found in the input file.
Wrong parameter dimension	variable has wrong dimension, expected: 13, actual: 3
Getter before evaluate()	Must call evaluate() before using getters
Unknown point name	Point NOPE not found
Unknown frame name	Frame NOPE not found

4 Typical Workflow and Examples

All code listings in this section are available as standalone Python scripts in the `examples/` directory alongside this document.

The standard usage pattern follows a *load – set – evaluate – get* cycle:

1. **Load.** Create an `Evaluator` from an LGC input file. This parses the file, builds the internal data structures, and sets the provisional parameter values.
2. **Set parameters** (optional). Replace the current parameter vector with custom values using `setParameters()`.
3. **Evaluate.** Call `evaluate()` to linearise the functional model at the current parameters. This computes the matrices A , P , A_2 and the misclosure vectors W , W_2 .
4. **Get results.** Retrieve any combination of matrices and vectors via the getter methods.

5. **Repeat** steps 2–4 as needed (e.g. for iterative solving or sensitivity analysis).

Alternatively, `solve()` performs the full iterative solution internally, returning the converged parameter vector directly.

Listing 1: Basic workflow example (`examples/01_basic_workflow.py`).

```
1 """
2 Basic workflow: Load a network, inspect it, and run a Least Squares
   Adjustment.
3 This script demonstrates the core API mechanics of pyLGC for
   surveyors.
4 """
5
6 import pyLGC
7
8 # 1. LOAD THE NETWORK
9 # The Evaluator parses the .lgc2 file and initializes the network
   geometry.
10 ev = pyLGC.Evaluator("Title-Example.lgc2")
11
12 # getProblemDimensions() returns an object containing the number of
   Observations (OIndex),
13 # Equations (EIndex), and Unknown parameters to estimate (UIndex).
14 idx = ev.getProblemDimensions()
15 print("=== Network Adjustment Summary ===")
16 print(f"Observations: {idx.OIndex} | Unknowns: {idx.UIndex}")
17
18
19 # 2. INSPECT THE STATIONS AND FRAMES
20 # getPoint() and getFrame() return objects representing network
   elements.
21 print("\n=== Inspecting Network Elements ===")
22 for name in ["H4.XBPF.22716.E", "H4.XBPF.22716.S"]:
23     pt = ev.getPoint(name)
24     # getRelativeUnknIndices() returns a NumPy array of indices.
25     # If the array is empty, the point is fully fixed. Otherwise, it
       is estimated.
26     status = "Estimated" if len(pt.getRelativeUnknIndices()) > 0 else
       "Fixed"
27     print(f"Station '{pt.getName()}' : {status}")
28
29
30 # 3. PRE-ADJUSTMENT EVALUATION
31 # IMPORTANT: You must call evaluate() before retrieving matrices or
32 # misclosures. This triggers the internal C++ engine to calculate
   theoretical
33 # measurements based on the current coordinates.
34 ev.evaluate()
35
36 # getMisclosure() returns a 1D NumPy array representing the
   difference between
37 # field measurements and theoretical measurements (the 'w' vector).
38 w_vector = ev.getMisclosure()
39 rms_before = (sum(x*x for x in w_vector) / len(w_vector))**0.5
40 print(f"\nInitial RMS of Misclosures: {rms_before:.6e}")
41
42
43 # 4. RUN THE LEAST SQUARES ADJUSTMENT
```

```

44 # solve() returns a tuple: (success_boolean, solution_array)
45 # The solution array contains the FINAL absolute estimated parameters
    (coordinates),
46 # NOT the iterative shifts (dX).
47 ok, solution = ev.solve()
48
49 if ok:
50     # 5. POST-ADJUSTMENT ANALYSIS
51     # setParameters() loads the final adjusted coordinates into
        memory, BUT...
52     ev.setParameters(solution)
53
54     # ...you MUST call evaluate() again to force the engine to
55     # re-calculate the network geometry with the new coordinates!
56     ev.evaluate()
57
58     # Now getMisclosure() returns the a posteriori residuals (the 'v'
        vector).
59     v_vector = ev.getMisclosure()
60
61     rms_after = (sum(x*x for x in v_vector) / len(v_vector))*0.5
62     print(f"Adjustment Succeeded! Post-Adjustment RMS: {rms_after:.6e
        }")
63 else:
64     print("Adjustment Failed to Converge!")

```

4.1 Sparse Matrix Usage with SciPy

All sparse matrices are returned in COO (coordinate) format as plain arrays, compatible with `scipy.sparse`:

Listing 2: Converting COO tuples to SciPy sparse matrices (examples/02_sparse_matrices.py).

```

1  """Converting COO tuples to SciPy sparse matrices."""
2
3  from scipy.sparse import coo_matrix
4  import pyLGC
5
6  ev = pyLGC.Evaluator("Title-Example.lgc2")
7  ev.evaluateAtParameters()
8
9  # Convert COO tuple to scipy sparse matrix
10 rows, cols, vals, nr, nc = ev.getFirstDesignMatrix()
11 A = coo_matrix((vals, (rows, cols)), shape=(nr, nc)).tocsr()
12
13 rows, cols, vals, nr, nc = ev.getWeightMatrix()
14 P = coo_matrix((vals, (rows, cols)), shape=(nr, nc)).tocsr()
15
16 w = ev.getMisclosure()

```

4.2 Building the Normal Equations in SciPy

Given the matrices from the evaluator, the normal equation (14) can be assembled and solved in Python:

Listing 3: Forming and solving the normal equations (`examples/03_normal_equations.py`).

```

1  """Forming and solving the normal equations."""
2
3  from scipy.sparse import coo_matrix
4  from scipy.sparse.linalg import spsolve
5  import pyLGC
6
7  ev = pyLGC.Evaluator("Title-Example.lgc2")
8  ev.evaluateAtParameters()
9
10 # Build sparse matrices
11 def to_csr(tup):
12     rows, cols, vals, nr, nc = tup
13     return coo_matrix((vals, (rows, cols)), shape=(nr, nc)).tocsr()
14
15 A = to_csr(ev.getFirstDesignMatrix())
16 P = to_csr(ev.getWeightMatrix())
17 w = ev.getMisclosure()
18
19 # Normal equations:  $(A^T P A) dx = -A^T P w$ 
20 N = A.T @ P @ A
21 n = A.T @ P @ w
22 dx = spsolve(N, -n)

```

4.3 Gauss–Newton Iteration

The following example implements a simple Gauss–Newton solver using the matrices provided by the evaluator. At each iteration the normal equations (14) are formed and solved, and the parameter vector is updated until the step size falls below a tolerance.

Listing 4: Simple Gauss–Newton solver, validated against the built-in solver (`examples/04_gauss_newton.py`).

```

1  """Simple Gauss-Newton solver, validated against the built-in solver.
2      """
3
4  import numpy as np
5  from numpy.linalg import norm
6  from scipy.sparse import coo_matrix
7  from scipy.sparse.linalg import spsolve
8  import pyLGC
9
10 def to_csr(tup):
11     rows, cols, vals, nr, nc = tup
12     return coo_matrix((vals, (rows, cols)), shape=(nr, nc)).tocsr()
13
14 ev = pyLGC.Evaluator("Title-Example.lgc2")
15
16 # 1. Reference solution from the built-in solver
17 x0 = ev.getEstimatedParameters().copy()
18 ok, x_ref = ev.solve()
19
20 # 2. Reset to provisional values and run own Gauss-Newton
21 x = x0.copy()
22 for k in range(50):
23     ev.setParameters(x)
24     ev.evaluateAtParameters()

```

```

24
25     A = to_csr(ev.getFirstDesignMatrix())
26     P = to_csr(ev.getWeightMatrix())
27     w = ev.getMisclosure()
28
29     # Solve normal equations: (A^T P A) dx = -A^T P w
30     N = A.T @ P @ A
31     dx = spsolve(N, -A.T @ P @ w)
32     x = x + dx
33
34     print(f"iter {k:2d}   |dx| = {norm(dx):.2e}   |w| = {norm(w):.2e}")
35     if norm(dx, np.inf) < 1e-8:
36         print("Converged.")
37         break
38
39     # 3. Compare the two solutions
40     print(f"||x_gn - x_ref|| = {norm(x - x_ref):.2e}")

```

References

- [1] J. Gutekunst, *LGC – Position Estimation Software: Principal Concepts and Methods*, CERN, 2024. <https://edms.cern.ch/document/3067818/1>
- [2] M. Barbier, *Least Square Equations for LGC2*, CERN, 2015. <https://edms.cern.ch/document/1475972/2>
- [3] D. Wells and E. Krakiwsky, *The Method of Least Squares*, University of New Brunswick, Department of Geodesy and Geomatics Engineering, Lecture Notes No. 18, 1971.